

House Prices : Linear Regression

```
In [24]: import pandas as pd
import numpy as np
from sklearn.model_selection import KFold
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
import matplotlib.pyplot as plt
import seaborn as sns
```

1. Reading the dataset into a Pandas DataFrame.

```
In [25]: df = pd.read_csv('train.csv', keep_default_na=False)
df.head()
```

```
Out[25]:
```

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandConto
0	1	60	RL	65	8450	Pave	NA	Reg	L
1	2	20	RL	80	9600	Pave	NA	Reg	L
2	3	60	RL	68	11250	Pave	NA	IR1	L
3	4	70	RL	60	9550	Pave	NA	IR1	L
4	5	60	RL	84	14260	Pave	NA	IR1	L

5 rows × 81 columns



```
In [26]: df.tail()
```

```
Out[26]:
```

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	Land
1455	1456	60	RL	62	7917	Pave	NA	Reg	
1456	1457	20	RL	85	13175	Pave	NA	Reg	
1457	1458	70	RL	66	9042	Pave	NA	Reg	
1458	1459	20	RL	68	9717	Pave	NA	Reg	
1459	1460	20	RL	75	9937	Pave	NA	Reg	

5 rows × 81 columns



2 Data Preprocessing

2.1 Data Types and non-null values

```
In [27]: df.info()
```

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 1460 entries, 0 to 1459

Data columns (total 81 columns):

#	Column	Non-Null Count	Dtype
0	Id	1460 non-null	int64
1	MSSubClass	1460 non-null	int64
2	MSZoning	1460 non-null	object
3	LotFrontage	1460 non-null	object
4	LotArea	1460 non-null	int64
5	Street	1460 non-null	object
6	Alley	1460 non-null	object
7	LotShape	1460 non-null	object
8	LandContour	1460 non-null	object
9	Utilities	1460 non-null	object
10	LotConfig	1460 non-null	object
11	LandSlope	1460 non-null	object
12	Neighborhood	1460 non-null	object
13	Condition1	1460 non-null	object
14	Condition2	1460 non-null	object
15	BldgType	1460 non-null	object
16	HouseStyle	1460 non-null	object
17	OverallQual	1460 non-null	int64
18	OverallCond	1460 non-null	int64
19	YearBuilt	1460 non-null	int64
20	YearRemodAdd	1460 non-null	int64
21	RoofStyle	1460 non-null	object
22	RoofMatl	1460 non-null	object
23	Exterior1st	1460 non-null	object
24	Exterior2nd	1460 non-null	object
25	MasVnrType	1460 non-null	object
26	MasVnrArea	1460 non-null	object
27	ExterQual	1460 non-null	object
28	ExterCond	1460 non-null	object
29	Foundation	1460 non-null	object
30	BsmtQual	1460 non-null	object
31	BsmtCond	1460 non-null	object
32	BsmtExposure	1460 non-null	object
33	BsmtFinType1	1460 non-null	object
34	BsmtFinSF1	1460 non-null	int64
35	BsmtFinType2	1460 non-null	object
36	BsmtFinSF2	1460 non-null	int64
37	BsmtUnfSF	1460 non-null	int64
38	TotalBsmtSF	1460 non-null	int64
39	Heating	1460 non-null	object
40	HeatingQC	1460 non-null	object
41	CentralAir	1460 non-null	object
42	Electrical	1460 non-null	object
43	1stFlrSF	1460 non-null	int64
44	2ndFlrSF	1460 non-null	int64
45	LowQualFinSF	1460 non-null	int64
46	GrLivArea	1460 non-null	int64
47	BsmtFullBath	1460 non-null	int64
48	BsmtHalfBath	1460 non-null	int64
49	FullBath	1460 non-null	int64
50	HalfBath	1460 non-null	int64

51	BedroomAbvGr	1460	non-null	int64
52	KitchenAbvGr	1460	non-null	int64
53	KitchenQual	1460	non-null	object
54	TotRmsAbvGrd	1460	non-null	int64
55	Functional	1460	non-null	object
56	Fireplaces	1460	non-null	int64
57	FireplaceQu	1460	non-null	object
58	GarageType	1460	non-null	object
59	GarageYrBlt	1460	non-null	object
60	GarageFinish	1460	non-null	object
61	GarageCars	1460	non-null	int64
62	GarageArea	1460	non-null	int64
63	GarageQual	1460	non-null	object
64	GarageCond	1460	non-null	object
65	PavedDrive	1460	non-null	object
66	WoodDeckSF	1460	non-null	int64
67	OpenPorchSF	1460	non-null	int64
68	EnclosedPorch	1460	non-null	int64
69	3SsnPorch	1460	non-null	int64
70	ScreenPorch	1460	non-null	int64
71	PoolArea	1460	non-null	int64
72	PoolQC	1460	non-null	object
73	Fence	1460	non-null	object
74	MiscFeature	1460	non-null	object
75	MiscVal	1460	non-null	int64
76	MoSold	1460	non-null	int64
77	YrSold	1460	non-null	int64
78	SaleType	1460	non-null	object
79	SaleCondition	1460	non-null	object
80	SalePrice	1460	non-null	int64

dtypes: int64(35), object(46)

memory usage: 924.0+ KB

2.2 Statistical Summary

In [28]: `df.describe()`


Out[28]:

	Id	MSSubClass	LotArea	OverallQual	OverallCond	YearBuilt
count	1460.000000	1460.000000	1460.000000	1460.000000	1460.000000	1460.000000
mean	730.500000	56.897260	10516.828082	6.099315	5.575342	1971.267808
std	421.610009	42.300571	9981.264932	1.382997	1.112799	30.202904
min	1.000000	20.000000	1300.000000	1.000000	1.000000	1872.000000
25%	365.750000	20.000000	7553.500000	5.000000	5.000000	1954.000000
50%	730.500000	50.000000	9478.500000	6.000000	5.000000	1973.000000
75%	1095.250000	70.000000	11601.500000	7.000000	6.000000	2000.000000
max	1460.000000	190.000000	215245.000000	10.000000	9.000000	2010.000000

8 rows × 35 columns



```
In [29]: # Plotting the distribution of SalePrice
plt.figure(figsize=(10, 6))
plt.hist(df['SalePrice'], bins=30, edgecolor='black')
plt.title('Distribution of SalePrice')
plt.xlabel('SalePrice')
plt.ylabel('Frequency')
plt.grid(axis='y', alpha=0.75)
plt.show()
```



```
In [30]: numerical_columns = df.select_dtypes(include=['int64', 'float64']).columns
print("Numerical Columns:", numerical_columns)
```

```
Numerical Columns: ['Id', 'MSSubClass', 'LotArea', 'OverallQual', 'OverallCo
nd', 'YearBuilt', 'YearRemodAdd', 'BsmtFinSF1', 'BsmtFinSF2', 'BsmtUnfSF',
'TotalBsmtSF', '1stFlrSF', '2ndFlrSF', 'LowQualFinSF', 'GrLivArea', 'BsmtFul
lBath', 'BsmtHalfBath', 'FullBath', 'HalfBath', 'BedroomAbvGr', 'KitchenAbvG
r', 'TotRmsAbvGrd', 'Fireplaces', 'GarageCars', 'GarageArea', 'WoodDeckSF',
'OpenPorchSF', 'EnclosedPorch', '3SsnPorch', 'ScreenPorch', 'PoolArea', 'Mis
cVal', 'MoSold', 'YrSold', 'SalePrice']
```

```
In [31]: #List of Categorical variables
categorical_vars = df.select_dtypes(include=['object']).columns.tolist()
print("Categorical Variables:", categorical_vars)
#LotFrontage is a numerical variable, which contains 'NA' replacing by mean
df['LotFrontage'] = df['LotFrontage'].replace('NA', np.nan)
df['LotFrontage'] = df['LotFrontage'].astype(float)
df['LotFrontage'].fillna(df['LotFrontage'].mean(), inplace=True)
df.head()
```

```
Categorical Variables: ['MSZoning', 'LotFrontage', 'Street', 'Alley', 'LotSh
ape', 'LandContour', 'Utilities', 'LotConfig', 'LandSlope', 'Neighborhood',
'Condition1', 'Condition2', 'BldgType', 'HouseStyle', 'RoofStyle', 'RoofMat
l', 'Exterior1st', 'Exterior2nd', 'MasVnrType', 'MasVnrArea', 'ExterQual',
'ExterCond', 'Foundation', 'BsmtQual', 'BsmtCond', 'BsmtExposure', 'BsmtFinT
ype1', 'BsmtFinType2', 'Heating', 'HeatingQC', 'CentralAir', 'Electrical',
'KitchenQual', 'Functional', 'FireplaceQu', 'GarageType', 'GarageYrBlt', 'Ga
rageFinish', 'GarageQual', 'GarageCond', 'PavedDrive', 'PoolQC', 'Fence', 'M
iscFeature', 'SaleType', 'SaleCondition']
```

/tmp/ipykernel_11894/2609515672.py:7: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.


For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or 'df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df['LotFrontage'].fillna(df['LotFrontage'].mean(), inplace=True)
```

```
Out[31]:
```

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandConto
0	1	60	RL	65.0	8450	Pave	NA	Reg	L
1	2	20	RL	80.0	9600	Pave	NA	Reg	L
2	3	60	RL	68.0	11250	Pave	NA	IR1	L
3	4	70	RL	60.0	9550	Pave	NA	IR1	L
4	5	60	RL	84.0	14260	Pave	NA	IR1	L

5 rows × 81 columns



2.3 Handling Missing Values

```
In [32]: df.isnull().sum()
```

```
Out[32]: Id                0
         MSSubClass        0
         MSZoning          0
         LotFrontage       0
         LotArea           0
         ..
         MoSold            0
         YrSold            0
         SaleType          0
         SaleCondition     0
         SalePrice         0
         Length: 81, dtype: int64
```

3. Splitting the dataset

Divide the dataset into two sets using k-fold cross validation technique entitled to train and test set respectively.

```
In [33]: # First, let's prepare our features and handle any remaining missing values
# Select relevant numerical features for our models
features_model1 = ['LotFrontage', 'LotArea']
features_model2 = ['LotFrontage', 'LotArea', 'OverallQual', 'OverallCond']
features_model3 = ['LotFrontage', 'LotArea', 'OverallQual', 'OverallCond', '

# Check for missing values in our target variable and features
print("Missing values in SalePrice:", df['SalePrice'].isnull().sum())
print("Missing values in features:")
for feature in features_model3:
    missing_count = df[feature].isnull().sum()
    print(f"{feature}: {missing_count}")

# Check data types
print(f"\nData types for features:")
for feature in features_model3:
    print(f"{feature}: {df[feature].dtype}")

# Remove any rows with missing values in our target variable or key features
print(f"\nOriginal dataset shape: {df.shape}")
df_clean = df.dropna(subset=['SalePrice'] + features_model3)
print(f"Dataset shape after removing missing values: {df_clean.shape}")

# Prepare target variable
y = df_clean['SalePrice']

# Prepare feature sets
X_model1 = df_clean[features_model1]
X_model2 = df_clean[features_model2]
X_model3 = df_clean[features_model3]

print(f"\nFeature sets prepared:")
print(f"Model 1 features: {features_model1}")
```

```

print(f"Model 2 features: {features_model2}")
print(f"Model 3 features: {features_model3}")
print(f"\nTarget variable (y) shape: {y.shape}")
print(f"Model 1 features shape: {X_model1.shape}")
print(f"Model 2 features shape: {X_model2.shape}")
print(f"Model 3 features shape: {X_model3.shape}")

```

Missing values in SalePrice: 0

Missing values in features:

LotFrontage: 0

LotArea: 0

OverallQual: 0

OverallCond: 0

1stFlrSF: 0

GrLivArea: 0

Data types for features:

LotFrontage: float64

LotArea: int64

OverallQual: int64

OverallCond: int64

1stFlrSF: int64

GrLivArea: int64

Original dataset shape: (1460, 81)

Dataset shape after removing missing values: (1460, 81)

Feature sets prepared:

Model 1 features: ['LotFrontage', 'LotArea']

Model 2 features: ['LotFrontage', 'LotArea', 'OverallQual', 'OverallCond']

Model 3 features: ['LotFrontage', 'LotArea', 'OverallQual', 'OverallCond',
'1stFlrSF', 'GrLivArea']

Target variable (y) shape: (1460,)

Model 1 features shape: (1460, 2)

Model 2 features shape: (1460, 4)

Model 3 features shape: (1460, 6)

```

In [34]: k = 5
kf = KFold(n_splits=k, shuffle=True, random_state=42)

# Initialize lists to store results for each model
models_data = {
    'Simple Linear Regression (LotArea only)': {'features': ['LotArea'], 'X': X_model1, 'y': y},
    'Model 1 (LotFrontage, LotArea)': {'features': features_model1, 'X': X_model1, 'y': y},
    'Model 2 (LotFrontage, LotArea, OverallQual, OverallCond)': {'features': features_model2, 'X': X_model2, 'y': y},
    'Model 3 (All features)': {'features': features_model3, 'X': X_model3, 'y': y}
}

# Store results for each model
results = {'model_name': {
    'train_mse': [], 'test_mse': [], 'train_r2': [], 'test_r2': [],
    'coefficients': [], 'intercepts': []
} for model_name in models_data.keys()}

print("Starting K-Fold Cross Validation...")

```



```
print(f"Using {k}-fold cross validation")
print("="*60)
```

Starting K-Fold Cross Validation...

Using 5-fold cross validation

=====

```
In [35]: fold = 1
for train_index, test_index in kf.split(y):
    print(f"\nFold {fold}:")
    print("-" * 20)

    # Split data for this fold
    y_train, y_test = y.iloc[train_index], y.iloc[test_index]

    # Train and evaluate each model
    for model_name, model_info in models_data.items():
        X = model_info['X']
        X_train, X_test = X.iloc[train_index], X.iloc[test_index]

        # Create and train linear regression model
        lr = LinearRegression()
        lr.fit(X_train, y_train)

        # Make predictions
        y_train_pred = lr.predict(X_train)
        y_test_pred = lr.predict(X_test)

        # Calculate metrics
        train_mse = mean_squared_error(y_train, y_train_pred)
        test_mse = mean_squared_error(y_test, y_test_pred)
        train_r2 = r2_score(y_train, y_train_pred)
        test_r2 = r2_score(y_test, y_test_pred)

        # Store results
        results[model_name]['train_mse'].append(train_mse)
        results[model_name]['test_mse'].append(test_mse)
        results[model_name]['train_r2'].append(train_r2)
        results[model_name]['test_r2'].append(test_r2)
        results[model_name]['coefficients'].append(lr.coef_)
        results[model_name]['intercepts'].append(lr.intercept_)

        print(f"{model_name}:")
        print(f"  Train MSE: {train_mse:.2f}, Test MSE: {test_mse:.2f}")
        print(f"  Train R²: {train_r2:.4f}, Test R²: {test_r2:.4f}")

    fold += 1

print("\nK-Fold Cross Validation completed!")
print("="*60)
```

Fold 1:

Simple Linear Regression (LotArea only):

Train MSE: 5541876621.93, Test MSE: 7189094014.83

Train R^2 : 0.0709, Test R^2 : 0.0627

Model 1 (LotFrontage, LotArea):

Train MSE: 5187706359.35, Test MSE: 6391769244.79

Train R^2 : 0.1302, Test R^2 : 0.1667

Model 2 (LotFrontage, LotArea, OverallQual, OverallCond):

Train MSE: 2031959572.87, Test MSE: 2326923040.59

Train R^2 : 0.6593, Test R^2 : 0.6966

Model 3 (All features):

Train MSE: 1594275967.62, Test MSE: 1702504242.03

Train R^2 : 0.7327, Test R^2 : 0.7780

Fold 2:

Simple Linear Regression (LotArea only):

Train MSE: 5743207806.18, Test MSE: 6393512739.93

Train R^2 : 0.0700, Test R^2 : 0.0597

Model 1 (LotFrontage, LotArea):

Train MSE: 5300597764.38, Test MSE: 5910512870.46

Train R^2 : 0.1417, Test R^2 : 0.1307

Model 2 (LotFrontage, LotArea, OverallQual, OverallCond):

Train MSE: 2053062935.08, Test MSE: 2218379551.32

Train R^2 : 0.6676, Test R^2 : 0.6737

Model 3 (All features):

Train MSE: 1657148893.40, Test MSE: 1443066802.25

Train R^2 : 0.7317, Test R^2 : 0.7878

Fold 3:

Simple Linear Regression (LotArea only):

Train MSE: 6044530616.25, Test MSE: 5181360177.35

Train R^2 : 0.0704, Test R^2 : 0.0621

Model 1 (LotFrontage, LotArea):

Train MSE: 5505585780.80, Test MSE: 5189460880.98

Train R^2 : 0.1533, Test R^2 : 0.0607

Model 2 (LotFrontage, LotArea, OverallQual, OverallCond):

Train MSE: 1998877019.06, Test MSE: 2545077057.47

Train R^2 : 0.6926, Test R^2 : 0.5393

Model 3 (All features):

Train MSE: 1312417880.99, Test MSE: 3069513195.16

Train R^2 : 0.7982, Test R^2 : 0.4444

Fold 4:

Simple Linear Regression (LotArea only):

Train MSE: 5790195671.80, Test MSE: 6500497232.62

Train R^2 : 0.0828, Test R^2 : -0.0353

Model 1 (LotFrontage, LotArea):

Train MSE: 5432144092.74, Test MSE: 5564972429.36

Train R^2 : 0.1395, Test R^2 : 0.1137

Model 2 (LotFrontage, LotArea, OverallQual, OverallCond):

Train MSE: 2121618371.57, Test MSE: 2002857872.51

Train R^2 : 0.6639, Test R^2 : 0.6810

Model 3 (All features):

Train MSE: 1676454936.75, Test MSE: 1351667140.21

Train R²: 0.7344, Test R²: 0.7847

Fold 5:

Simple Linear Regression (LotArea only):

Train MSE: 6160516422.26, Test MSE: 4714667164.94

Train R²: 0.0633, Test R²: 0.0980

Model 1 (LotFrontage, LotArea):

Train MSE: 5601182550.95, Test MSE: 4722809022.10

Train R²: 0.1483, Test R²: 0.0964

Model 2 (LotFrontage, LotArea, OverallQual, OverallCond):

Train MSE: 2168161716.58, Test MSE: 1756509217.54

Train R²: 0.6703, Test R²: 0.6639

Model 3 (All features):

Train MSE: 1737070127.12, Test MSE: 1107656712.95

Train R²: 0.7359, Test R²: 0.7881

K-Fold Cross Validation completed!

=====

```
In [36]: print("\n" + "="*80)
print("SUMMARY OF MODEL PERFORMANCE ACROSS ALL FOLDS")
print("="*80)

summary_results = {}
for model_name in models_data.keys():
    avg_train_mse = np.mean(results[model_name]['train_mse'])
    avg_test_mse = np.mean(results[model_name]['test_mse'])
    avg_train_r2 = np.mean(results[model_name]['train_r2'])
    avg_test_r2 = np.mean(results[model_name]['test_r2'])

    std_train_mse = np.std(results[model_name]['train_mse'])
    std_test_mse = np.std(results[model_name]['test_mse'])
    std_train_r2 = np.std(results[model_name]['train_r2'])
    std_test_r2 = np.std(results[model_name]['test_r2'])

    summary_results[model_name] = {
        'avg_train_mse': avg_train_mse, 'avg_test_mse': avg_test_mse,
        'avg_train_r2': avg_train_r2, 'avg_test_r2': avg_test_r2,
        'std_train_mse': std_train_mse, 'std_test_mse': std_test_mse,
        'std_train_r2': std_train_r2, 'std_test_r2': std_test_r2
    }

    print(f"\n{model_name}:")
    print(f"  Average Train MSE: {avg_train_mse:.2f} (±{std_train_mse:.2f})")
    print(f"  Average Test MSE: {avg_test_mse:.2f} (±{std_test_mse:.2f})")
    print(f"  Average Train R²: {avg_train_r2:.4f} (±{std_train_r2:.4f})")
    print(f"  Average Test R²: {avg_test_r2:.4f} (±{std_test_r2:.4f})")

print("\n" + "="*80)
```

```
=====
=====
SUMMARY OF MODEL PERFORMANCE ACROSS ALL FOLDS
=====
=====
```

Simple Linear Regression (LotArea only):

Average Train MSE: 5856065427.68 (± 220907133.67)
 Average Test MSE: 5995826265.93 (± 910105366.24)
 Average Train R^2 : 0.0715 (± 0.0063)
 Average Test R^2 : 0.0495 (± 0.0447)

Model 1 (LotFrontage, LotArea):

Average Train MSE: 5405443309.64 (± 146599634.94)
 Average Test MSE: 5555904889.54 (± 574948379.69)
 Average Train R^2 : 0.1426 (± 0.0079)
 Average Test R^2 : 0.1136 (± 0.0352)

Model 2 (LotFrontage, LotArea, OverallQual, OverallCond):

Average Train MSE: 2074735923.03 (± 61612418.77)
 Average Test MSE: 2169949347.89 (± 270761639.02)
 Average Train R^2 : 0.6707 (± 0.0115)
 Average Test R^2 : 0.6509 (± 0.0568)

Model 3 (All features):

Average Train MSE: 1595473561.18 (± 148683132.60)
 Average Test MSE: 1734881618.52 (± 693932279.90)
 Average Train R^2 : 0.7466 (± 0.0258)
 Average Test R^2 : 0.7166 (± 0.1362)

```
=====
=====

In [37]: print("DETAILED COEFFICIENT ANALYSIS")
          print("="*80)

          for model_name, model_info in models_data.items():
              print(f"\n{model_name}:")
              print("-" * len(model_name))

              # Calculate average coefficients and intercept
              avg_coefficients = np.mean(results[model_name]['coefficients'], axis=0)
              avg_intercept = np.mean(results[model_name]['intercepts'])

              std_coefficients = np.std(results[model_name]['coefficients'], axis=0)
              std_intercept = np.std(results[model_name]['intercepts'])

              print(f" Intercept: {avg_intercept:.2f} ( $\pm$ {std_intercept:.2f})")

              features = model_info['features']
              for i, feature in enumerate(features):
                  print(f" {feature}: {avg_coefficients[i]:.2f} ( $\pm$ {std_coefficients[i]:.2f})")

              # Display the regression equation
              equation = f"SalePrice = {avg_intercept:.2f}"
              for i, feature in enumerate(features):
```

```

        if avg_coefficients[i] >= 0:
            equation += f" + {avg_coefficients[i]:.2f}×{feature}"
        else:
            equation += f" - {abs(avg_coefficients[i]):.2f}×{feature}"

    print(f"\n Regression Equation: {equation}")

print("\n" + "="*80)

```

DETAILED COEFFICIENT ANALYSIS

```

=====
=====

```

Simple Linear Regression (LotArea only):

```
-----
```

Intercept: 157978.16 (±3664.22)
LotArea: 2.18 (±0.39)

Regression Equation: SalePrice = 157978.16 + 2.18×LotArea

Model 1 (LotFrontage, LotArea):

```
-----
```

Intercept: 94361.69 (±7890.40)
LotFrontage: 1013.78 (±123.14)
LotArea: 1.48 (±0.32)

Regression Equation: SalePrice = 94361.69 + 1013.78×LotFrontage + 1.48×LotArea

Model 2 (LotFrontage, LotArea, OverallQual, OverallCond):

```
-----
```

Intercept: -122296.92 (±12246.28)
LotFrontage: 414.37 (±86.96)
LotArea: 1.24 (±0.19)
OverallQual: 42961.40 (±754.04)
OverallCond: -143.80 (±414.39)

Regression Equation: SalePrice = -122296.92 + 414.37×LotFrontage + 1.24×LotArea + 42961.40×OverallQual - 143.80×OverallCond

Model 3 (All features):

```
-----
```

Intercept: -129666.74 (±14067.21)
LotFrontage: 67.12 (±65.07)
LotArea: 0.66 (±0.06)
OverallQual: 31123.91 (±741.01)
OverallCond: 1385.18 (±695.95)
1stFlrSF: 34.87 (±4.77)
GrLivArea: 40.16 (±3.93)

Regression Equation: SalePrice = -129666.74 + 67.12×LotFrontage + 0.66×LotArea + 31123.91×OverallQual + 1385.18×OverallCond + 34.87×1stFlrSF + 40.16×GrLivArea

```

=====
=====

```

```

In [38]: fig, ((ax1, ax2), (ax3, ax4)) = plt.subplots(2, 2, figsize=(15, 12))

# Prepare data for plotting
model_names = list(models_data.keys())
train_mse_means = [summary_results[name]['avg_train_mse'] for name in model_names]
test_mse_means = [summary_results[name]['avg_test_mse'] for name in model_names]
train_r2_means = [summary_results[name]['avg_train_r2'] for name in model_names]
test_r2_means = [summary_results[name]['avg_test_r2'] for name in model_names]

train_mse_stds = [summary_results[name]['std_train_mse'] for name in model_names]
test_mse_stds = [summary_results[name]['std_test_mse'] for name in model_names]
train_r2_stds = [summary_results[name]['std_train_r2'] for name in model_names]
test_r2_stds = [summary_results[name]['std_test_r2'] for name in model_names]

# Shorten model names for better visualization
short_names = ['Simple LR', 'Model 1', 'Model 2', 'Model 3']

# Plot 1: Training MSE Comparison
ax1.bar(short_names, train_mse_means, yerr=train_mse_stds, capsize=5, alpha=0.7)
ax1.set_title('Training MSE Comparison')
ax1.set_ylabel('Mean Squared Error')
ax1.tick_params(axis='x', rotation=45)

# Plot 2: Test MSE Comparison
ax2.bar(short_names, test_mse_means, yerr=test_mse_stds, capsize=5, alpha=0.7)
ax2.set_title('Test MSE Comparison')
ax2.set_ylabel('Mean Squared Error')
ax2.tick_params(axis='x', rotation=45)

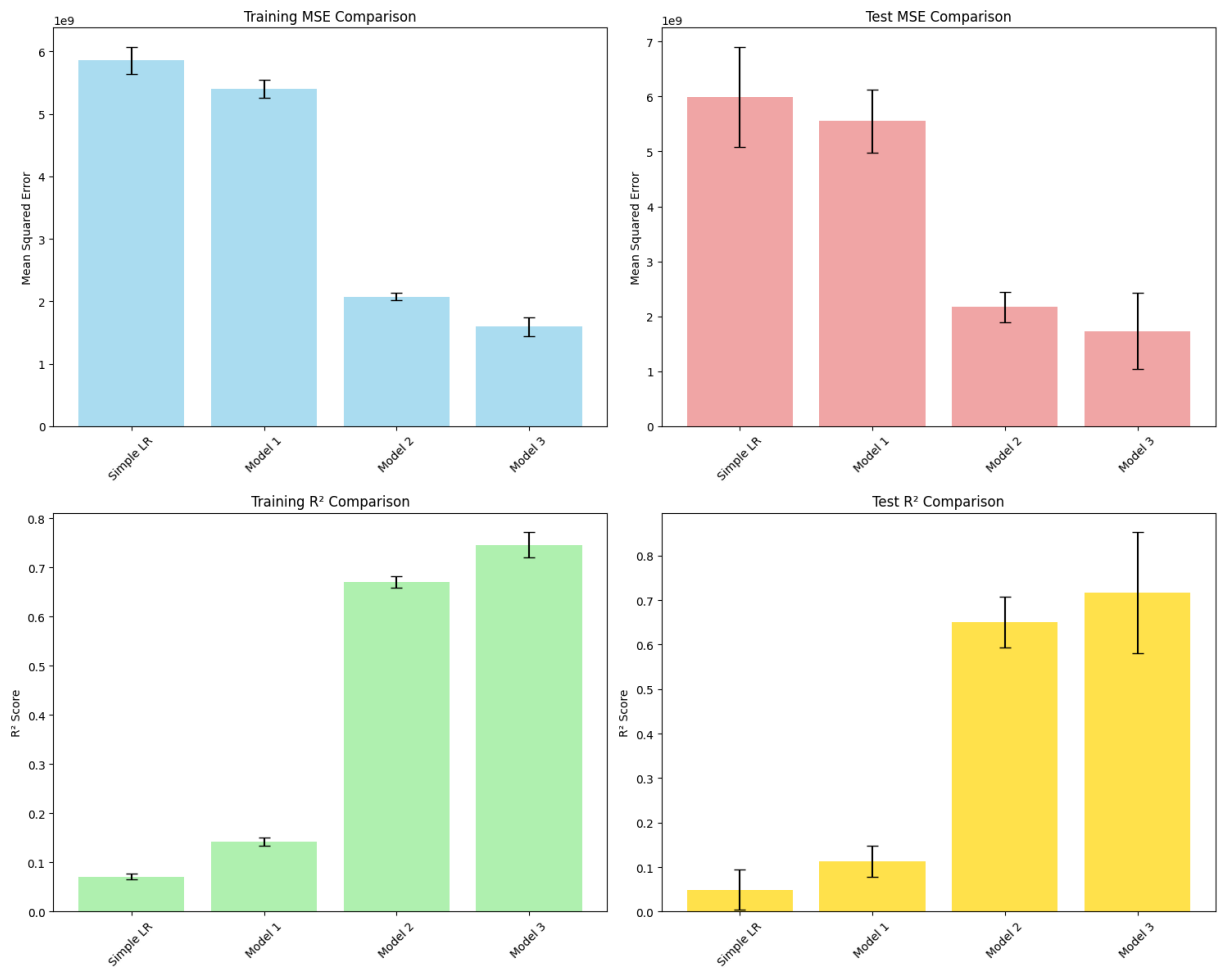
# Plot 3: Training R2 Comparison
ax3.bar(short_names, train_r2_means, yerr=train_r2_stds, capsize=5, alpha=0.7)
ax3.set_title('Training R2 Comparison')
ax3.set_ylabel('R2 Score')
ax3.tick_params(axis='x', rotation=45)

# Plot 4: Test R2 Comparison
ax4.bar(short_names, test_r2_means, yerr=test_r2_stds, capsize=5, alpha=0.7)
ax4.set_title('Test R2 Comparison')
ax4.set_ylabel('R2 Score')
ax4.tick_params(axis='x', rotation=45)

plt.tight_layout()
plt.show()

# Also create a summary table
print("\nMODEL PERFORMANCE SUMMARY TABLE")
print("="*80)
print(f"{'Model':<30} {'Train MSE':<12} {'Test MSE':<12} {'Train R2':<12} {'Test R2':<12}")
print("-" * 80)
for i, name in enumerate(model_names):
    short_name = short_names[i]
    print(f"{short_name:<30} {train_mse_means[i]:<12.2f} {test_mse_means[i]:<12.2f} {train_r2_means[i]:<12.2f} {test_r2_means[i]:<12.2f}")
print("="*80)

```



MODEL PERFORMANCE SUMMARY TABLE

Model	Train MSE	Test MSE	Train R ²	Test R ²
Simple LR	5856065427.68	5995826265.93	0.0715	0.04
Model 1	5405443309.64	5555904889.54	0.1426	0.11
Model 2	2074735923.03	2169949347.89	0.6707	0.65
Model 3	1595473561.18	1734881618.52	0.7466	0.71

```
In [39]: # Model Analysis and Conclusions
print("MODEL ANALYSIS AND INSIGHTS")
print("="*80)

# Find the best performing model
best_test_r2_idx = np.argmax(test_r2_means)
best_test_mse_idx = np.argmin(test_mse_means)

print(f"\nBest Model by Test R2 Score: {short_names[best_test_r2_idx]} (R2 =
```

```

print(f"Best Model by Test MSE: {short_names[best_test_mse_idx]} (MSE = {tes

# Calculate improvement from simple to complex models
simple_test_r2 = test_r2_means[0]
best_test_r2 = test_r2_means[best_test_r2_idx]
improvement = ((best_test_r2 - simple_test_r2) / simple_test_r2) * 100

print(f"\nImprovement from Simple Linear Regression to Best Model: {improven

# Check for overfitting
print(f"\nOverfitting Analysis (Train R2 - Test R2):")
for i, name in enumerate(short_names):
    overfitting = train_r2_means[i] - test_r2_means[i]
    print(f"    {name}: {overfitting:.4f}")

print(f"1. Simple Linear Regression (LotArea only) achieved R2 = {simple_tes

# Display feature importance (from the best model)
best_model_name = list(models_data.keys())[best_test_r2_idx]
best_features = models_data[best_model_name]['features']
best_coeffs = np.mean(results[best_model_name]['coefficients'], axis=0)

print(f"\nFeature Importance in Best Model ({short_names[best_test_r2_idx]})
feature_importance = list(zip(best_features, np.abs(best_coeffs)))
feature_importance.sort(key=lambda x: x[1], reverse=True)

for feature, importance in feature_importance:
    print(f"    {feature}: {importance:.2f}")

print("="*80)

```


MODEL ANALYSIS AND INSIGHTS

Best Model by Test R^2 Score: Model 3 ($R^2 = 0.7166$)
Best Model by Test MSE: Model 3 (MSE = 1734881618.52)

Improvement from Simple Linear Regression to Best Model: 1348.91%

Overfitting Analysis (Train R^2 - Test R^2):

Simple LR: 0.0220

Model 1: 0.0290

Model 2: 0.0198

Model 3: 0.0300

1. Simple Linear Regression (LotArea only) achieved $R^2 = 0.0495$

Feature Importance in Best Model (Model 3):

OverallQual: 31123.91

OverallCond: 1385.18

LotFrontage: 67.12

GrLivArea: 40.16

1stFlrSF: 34.87

LotArea: 0.66

4. Multiple Regression with Mixed Features (Numerical + Categorical)

```
In [40]: from sklearn.preprocessing import LabelEncoder, OneHotEncoder
import pandas as pd

print("Examining categorical features:")
print("Street unique values:", df['Street'].unique())
print("Street value counts:", df['Street'].value_counts())
print("\nNeighborhood unique values (first 10):", df['Neighborhood'].unique()
print("Neighborhood value counts (top 10):")
print(df['Neighborhood'].value_counts().head(10))

# Checking if we have YearBuilt or need to create Year feature
print(f"\nAvailable year-related columns:")
year_columns = [col for col in df.columns if 'year' in col.lower() or 'yr' in col.lower()]
print(year_columns)

# Using YearBuilt as our Year feature
if 'YearBuilt' in df.columns:
    print(f"Using YearBuilt as Year feature")
    print(f"YearBuilt range: {df['YearBuilt'].min()} - {df['YearBuilt'].max()}")
else:
    print("YearBuilt not found, will create a dummy Year feature")
```

```
Examining categorical features:
Street unique values: ['Pave' 'Grvl']
Street value counts: Street
Pave      1454
Grvl        6
Name: count, dtype: int64
```

```
Neighborhood unique values (first 10): ['CollgCr' 'Veenker' 'Crawfor' 'NoRid
ge' 'Mitchel' 'Somerst' 'NWAmes'
'OldTown' 'BrkSide' 'Sawyer']
Neighborhood value counts (top 10):
Neighborhood
NAMES      225
CollgCr    150
OldTown    113
Edwards    100
Somerst     86
Gilbert     79
NridgHt     77
Sawyer      74
NWAmes      73
SawyerW     59
Name: count, dtype: int64
```

```
Available year-related columns:
['YearBuilt', 'YearRemodAdd', 'GarageYrBlt', 'YrSold']
Using YearBuilt as Year feature
YearBuilt range: 1872 - 2010
```

```
In [41]: from sklearn.preprocessing import LabelEncoder

mixed_features_model4 = {
    'numerical': ['LotArea'],
    'categorical': ['Street']
}

mixed_features_model5 = {
    'numerical': ['LotArea', 'OverallCond'],
    'categorical': ['Street', 'Neighborhood']
}

mixed_features_model6 = {
    'numerical': ['LotArea', 'OverallCond', '1stFlrSF'],
    'categorical': ['Street', 'Neighborhood', 'YearBuilt'] # Using YearBuilt
}

def prepare_mixed_features_simple(df, numerical_features, categorical_features):
    """
    Prepare mixed features using Label Encoding for categorical variables
    to keep the exact number of features as specified in the assignment
    """
    # Start with numerical features
    X_mixed = df[numerical_features].copy()

    # Add label-encoded categorical features
    le = LabelEncoder()
```

```

    for cat_feature in categorical_features:
        if cat_feature == 'YearBuilt':
            # Use YearBuilt as-is (it's already numerical, but we'll treat it as categorical)
            X_mixed[f'{cat_feature}_encoded'] = le.fit_transform(df[cat_feature])
        else:
            # Standard label encoding for other categorical variables
            X_mixed[f'{cat_feature}_encoded'] = le.fit_transform(df[cat_feature])

    return X_mixed

print("Preparing mixed feature datasets with correct feature counts...")
print("Using Label Encoding to maintain exact feature counts as per assignment")

X_model4 = prepare_mixed_features_simple(df,
                                         mixed_features_model4['numerical'],
                                         mixed_features_model4['categorical'])

X_model5 = prepare_mixed_features_simple(df,
                                         mixed_features_model5['numerical'],
                                         mixed_features_model5['categorical'])

X_model6 = prepare_mixed_features_simple(df,
                                         mixed_features_model6['numerical'],
                                         mixed_features_model6['categorical'])

print(f"Model 4 shape: {X_model4.shape} (Expected: 2 features)")
print(f"Model 5 shape: {X_model5.shape} (Expected: 4 features)")
print(f"Model 6 shape: {X_model6.shape} (Expected: 6 features)")

print(f"\nModel 4 features: {list(X_model4.columns)}")
print(f"Model 5 features: {list(X_model5.columns)}")
print(f"Model 6 features: {list(X_model6.columns)}")

print(f"\nFeature specifications as per assignment:")
print(f"Model 4: LotArea, Street")
print(f"Model 5: LotArea, OverallCond, Street, Neighborhood")
print(f"Model 6: LotArea, OverallCond, Street, 1stFlrSF, Neighborhood, Year")

```

Preparing mixed feature datasets with correct feature counts...

Using Label Encoding to maintain exact feature counts as per assignment

Model 4 shape: (1460, 2) (Expected: 2 features)

Model 5 shape: (1460, 4) (Expected: 4 features)

Model 6 shape: (1460, 6) (Expected: 6 features)

Model 4 features: ['LotArea', 'Street_encoded']

Model 5 features: ['LotArea', 'OverallCond', 'Street_encoded', 'Neighborhood_encoded']

Model 6 features: ['LotArea', 'OverallCond', '1stFlrSF', 'Street_encoded', 'Neighborhood_encoded', 'YearBuilt_encoded']

Feature specifications as per assignment:

Model 4: LotArea, Street

Model 5: LotArea, OverallCond, Street, Neighborhood

Model 6: LotArea, OverallCond, Street, 1stFlrSF, Neighborhood, Year

```

In [42]: mixed_models_data = {
    'Model 4 (LotArea, Street)': {'X': X_model4, 'features': mixed_features_
    'Model 5 (LotArea, OverallCond, Street, Neighborhood)': {'X': X_model5,
    'Model 6 (LotArea, OverallCond, Street, 1stFlrSF, Neighborhood, Year)':
}

# Store results for mixed models
mixed_results = {model_name: {
    'train_mse': [], 'test_mse': [], 'train_r2': [], 'test_r2': [],
    'coefficients': [], 'intercepts': [], 'feature_names': []
} for model_name in mixed_models_data.keys()}

print("Starting K-Fold Cross Validation for Mixed Models...")
print(f"Using {k}-fold cross validation")
print("="*60)

# Perform K-Fold Cross Validation for mixed models
fold = 1
for train_index, test_index in kf.split(y):
    print(f"\nFold {fold}:")
    print("-" * 20)

    # Split data for this fold
    y_train, y_test = y.iloc[train_index], y.iloc[test_index]

    # Train and evaluate each mixed model
    for model_name, model_info in mixed_models_data.items():
        X = model_info['X']
        X_train, X_test = X.iloc[train_index], X.iloc[test_index]

        # Create and train linear regression model
        lr = LinearRegression()
        lr.fit(X_train, y_train)

        # Make predictions
        y_train_pred = lr.predict(X_train)
        y_test_pred = lr.predict(X_test)

        # Calculate metrics
        train_mse = mean_squared_error(y_train, y_train_pred)
        test_mse = mean_squared_error(y_test, y_test_pred)
        train_r2 = r2_score(y_train, y_train_pred)
        test_r2 = r2_score(y_test, y_test_pred)

        # Store results
        mixed_results[model_name]['train_mse'].append(train_mse)
        mixed_results[model_name]['test_mse'].append(test_mse)
        mixed_results[model_name]['train_r2'].append(train_r2)
        mixed_results[model_name]['test_r2'].append(test_r2)
        mixed_results[model_name]['coefficients'].append(lr.coef_)
        mixed_results[model_name]['intercepts'].append(lr.intercept_)
        mixed_results[model_name]['feature_names'] = list(X.columns)

    print(f"{model_name}:")
    print(f"  Train MSE: {train_mse:.2f}, Test MSE: {test_mse:.2f}")

```

```
        print(f"  Train R²: {train_r2:.4f}, Test R²: {test_r2:.4f}")

    fold += 1

print("\nK-Fold Cross Validation for Mixed Models completed!")
print("="*60)
```

Starting K-Fold Cross Validation for Mixed Models...
Using 5-fold cross validation

=====

Fold 1:

Model 4 (LotArea, Street):

Train MSE: 5501517875.25, Test MSE: 7066617147.62

Train R^2 : 0.0776, Test R^2 : 0.0787

Model 5 (LotArea, OverallCond, Street, Neighborhood):

Train MSE: 5255968054.67, Test MSE: 6627543389.01

Train R^2 : 0.1188, Test R^2 : 0.1360

Model 6 (LotArea, OverallCond, Street, 1stFlrSF, Neighborhood, Year):

Train MSE: 2731781241.83, Test MSE: 3112864268.77

Train R^2 : 0.5420, Test R^2 : 0.5942

Fold 2:

Model 4 (LotArea, Street):

Train MSE: 5681823026.86, Test MSE: 6353847486.19

Train R^2 : 0.0800, Test R^2 : 0.0655

Model 5 (LotArea, OverallCond, Street, Neighborhood):

Train MSE: 5420442699.55, Test MSE: 5973014669.19

Train R^2 : 0.1223, Test R^2 : 0.1215

Model 6 (LotArea, OverallCond, Street, 1stFlrSF, Neighborhood, Year):

Train MSE: 2770274250.89, Test MSE: 2928670148.33

Train R^2 : 0.5514, Test R^2 : 0.5693

Fold 3:

Model 4 (LotArea, Street):

Train MSE: 5963391668.40, Test MSE: 5251925666.86

Train R^2 : 0.0829, Test R^2 : 0.0494

Model 5 (LotArea, OverallCond, Street, Neighborhood):

Train MSE: 5683235215.79, Test MSE: 4919492989.28

Train R^2 : 0.1259, Test R^2 : 0.1095

Model 6 (LotArea, OverallCond, Street, 1stFlrSF, Neighborhood, Year):

Train MSE: 2633470532.62, Test MSE: 3653198890.14

Train R^2 : 0.5950, Test R^2 : 0.3387

Fold 4:

Model 4 (LotArea, Street):

Train MSE: 5756539473.50, Test MSE: 6312944192.43

Train R^2 : 0.0881, Test R^2 : -0.0054

Model 5 (LotArea, OverallCond, Street, Neighborhood):

Train MSE: 5430058237.50, Test MSE: 6261618030.70

Train R^2 : 0.1398, Test R^2 : 0.0028

Model 6 (LotArea, OverallCond, Street, 1stFlrSF, Neighborhood, Year):

Train MSE: 2857270497.66, Test MSE: 2644853107.39

Train R^2 : 0.5474, Test R^2 : 0.5788

Fold 5:

Model 4 (LotArea, Street):

Train MSE: 6093110687.71, Test MSE: 4703317777.04

Train R^2 : 0.0735, Test R^2 : 0.1002
 Model 5 (LotArea, OverallCond, Street, Neighborhood):
 Train MSE: 5772070001.46, Test MSE: 4560069474.27
 Train R^2 : 0.1223, Test R^2 : 0.1276
 Model 6 (LotArea, OverallCond, Street, 1stFlrSF, Neighborhood, Year):
 Train MSE: 2948996154.60, Test MSE: 2223596122.89
 Train R^2 : 0.5516, Test R^2 : 0.5746

K-Fold Cross Validation for Mixed Models completed!

=====

```
In [43]: # Calculate and display average performance metrics for mixed models
print("\n" + "="*80)
print("SUMMARY OF MIXED MODEL PERFORMANCE ACROSS ALL FOLDS")
print("="*80)

mixed_summary_results = {}
for model_name in mixed_models_data.keys():
    avg_train_mse = np.mean(mixed_results[model_name]['train_mse'])
    avg_test_mse = np.mean(mixed_results[model_name]['test_mse'])
    avg_train_r2 = np.mean(mixed_results[model_name]['train_r2'])
    avg_test_r2 = np.mean(mixed_results[model_name]['test_r2'])

    std_train_mse = np.std(mixed_results[model_name]['train_mse'])
    std_test_mse = np.std(mixed_results[model_name]['test_mse'])
    std_train_r2 = np.std(mixed_results[model_name]['train_r2'])
    std_test_r2 = np.std(mixed_results[model_name]['test_r2'])

    mixed_summary_results[model_name] = {
        'avg_train_mse': avg_train_mse, 'avg_test_mse': avg_test_mse,
        'avg_train_r2': avg_train_r2, 'avg_test_r2': avg_test_r2,
        'std_train_mse': std_train_mse, 'std_test_mse': std_test_mse,
        'std_train_r2': std_train_r2, 'std_test_r2': std_test_r2
    }

    print(f"\n{model_name}:")
    print(f"  Average Train MSE: {avg_train_mse:.2f} ( $\pm$ {std_train_mse:.2f})")
    print(f"  Average Test MSE: {avg_test_mse:.2f} ( $\pm$ {std_test_mse:.2f})")
    print(f"  Average Train  $R^2$ : {avg_train_r2:.4f} ( $\pm$ {std_train_r2:.4f})")
    print(f"  Average Test  $R^2$ : {avg_test_r2:.4f} ( $\pm$ {std_test_r2:.4f})")

print("\n" + "="*80)
```

```
=====
=====
SUMMARY OF MIXED MODEL PERFORMANCE ACROSS ALL FOLDS
=====
=====
```

Model 4 (LotArea, Street):

```
Average Train MSE: 5799276546.34 (±208592563.99)
Average Test MSE: 5937730454.03 (±846454578.12)
Average Train R²: 0.0804 (±0.0049)
Average Test R²: 0.0577 (±0.0357)
```

Model 5 (LotArea, OverallCond, Street, Neighborhood):

```
Average Train MSE: 5512354841.79 (±188471124.56)
Average Test MSE: 5668347710.49 (±794218108.66)
Average Train R²: 0.1258 (±0.0074)
Average Test R²: 0.0995 (±0.0491)
```

Model 6 (LotArea, OverallCond, Street, 1stFlrSF, Neighborhood, Year):

```
Average Train MSE: 2788358535.52 (±107769780.68)
Average Test MSE: 2912636507.51 (±476500887.22)
Average Train R²: 0.5575 (±0.0191)
Average Test R²: 0.5311 (±0.0965)
```

```
=====
=====
```

```
In [44]: # Extract LotArea coefficients from all models
print("LOTAREA COEFFICIENT ANALYSIS ACROSS ALL MODELS")
print("="*80)

# Function to extract LotArea coefficient from model results
def get_lotarea_coefficient(model_results, feature_names):
    """Extract the LotArea coefficient from model results"""
    coefficients = np.mean(model_results['coefficients'], axis=0)

    # Find the index of LotArea in the feature names
    if 'LotArea' in feature_names:
        lotarea_idx = feature_names.index('LotArea')
        return coefficients[lotarea_idx]
    else:
        return None

# Collect LotArea coefficients from all models
lotarea_coefficients = {}
model_labels = []

# Original models (1-3)
original_model_names = ['Simple Linear Regression (LotArea only)', 'Model 1 (LotArea, OverallQual, OverallCond)',
                        'Model 2 (LotFrontage, LotArea, OverallQual, OverallCond)', 'Model 3 (LotArea, OverallQual, OverallCond, Year)']

for model_name in original_model_names:
    if model_name in results:
        features = models_data[model_name]['features']
        coef = get_lotarea_coefficient(results[model_name], features)
        if coef is not None:
```



```

        lotarea_coefficients[model_name] = coef
        model_labels.append(model_name.split('(')[0].strip())

# Mixed models (4-6)
mixed_model_names = list(mixed_models_data.keys())
for model_name in mixed_model_names:
    feature_names = mixed_results[model_name]['feature_names']
    coef = get_lotarea_coefficient(mixed_results[model_name], feature_names)
    if coef is not None:
        lotarea_coefficients[model_name] = coef
        model_labels.append(model_name.split('(')[0].strip())

# Display the coefficients
print("LotArea coefficients across all models:")
print("-" * 50)
for i, (model_name, coef) in enumerate(lotarea_coefficients.items()):
    print(f"{model_labels[i]}: {coef:.4f}")

print("\n" + "="*80)

```

LOTAREA COEFFICIENT ANALYSIS ACROSS ALL MODELS

```

=====
====
LotArea coefficients across all models:
-----
Simple Linear Regression: 2.1849
Model 1: 1.4830
Model 2: 1.2357
Model 3: 0.6618
Model 4: 2.3371
Model 5: 2.2721
Model 6: 1.0959

=====
=====

```

```

In [45]: # Create visualization of LotArea coefficients across all models
plt.figure(figsize=(14, 8))

# Prepare data for plotting
coefficients_values = list(lotarea_coefficients.values())
short_model_names = ['Simple LR', 'Model 1', 'Model 2', 'Model 3', 'Model 4']

# Create the bar plot
bars = plt.bar(range(len(coefficients_values)), coefficients_values,
               color=['skyblue', 'lightcoral', 'lightgreen', 'gold', 'orange'],
               alpha=0.8, edgecolor='black', linewidth=1)

# Customize the plot
plt.title('LotArea Coefficient Comparison Across All Models', fontsize=16, fontweight='bold')
plt.xlabel('Models', fontsize=12, fontweight='bold')
plt.ylabel('LotArea Coefficient Value', fontsize=12, fontweight='bold')
plt.xticks(range(len(coefficients_values)), short_model_names, rotation=45)

# Add value labels on top of each bar
for i, (bar, value) in enumerate(zip(bars, coefficients_values)):

```

```

plt.text(bar.get_x() + bar.get_width()/2, bar.get_height() + max(coefficient, 0.1),
         f'{value:.3f}', ha='center', va='bottom', fontweight='bold', fontcolor='black')

# Add grid for better readability
plt.grid(axis='y', alpha=0.3, linestyle='--')

# Add a horizontal line at y=0 for reference
plt.axhline(y=0, color='red', linestyle='--', alpha=0.5)

plt.tight_layout()
plt.show()


# Create a detailed comparison table
print("\nDETAILED LOTAREA COEFFICIENT COMPARISON")
print("="*80)
print(f"{'Model':<15} {'LotArea Coeff':<15} {'Features Description':<50}")
print("-" * 80)

model_descriptions = [
    "LotArea only",
    "LotFrontage + LotArea",
    "LotFrontage + LotArea + OverallQual + OverallCond",
    "LotFrontage + LotArea + OverallQual + OverallCond + 1stFlrSF + GrLivArea",
    "LotArea + Street (categorical)",
    "LotArea + OverallCond + Street + Neighborhood (categorical)",
    "LotArea + OverallCond + Street + 1stFlrSF + Neighborhood + Year (categorical)"
]

for i, (coef, desc) in enumerate(zip(coefficients_values, model_descriptions)):
    print(f"{'short_model_names[i]':<15} {coef:<15.4f} {desc:<50}")

print("="*80)

```



DETAILED LOTAREA COEFFICIENT COMPARISON

Model	LotArea Coeff	Features Description
Simple LR	2.1849	LotArea only
Model 1	1.4830	LotFrontage + LotArea
Model 2	1.2357	LotFrontage + LotArea + OverallQual + OverallCond
Model 3	0.6618	LotFrontage + LotArea + OverallQual + OverallCond + 1stFlrSF + GrLivArea
Model 4	2.3371	LotArea + Street (categorical)
Model 5	2.2721	LotArea + OverallCond + Street + Neighborhood (categorical)
Model 6	1.0959	LotArea + OverallCond + Street + 1stFlrSF + Neighborhood + Year (categorical)

```
In [46]: # Analysis of LotArea coefficient changes
print("ANALYSIS OF LOTAREA COEFFICIENT CHANGES")
print("="*80)

# Calculate percentage changes
base_coef = coefficients_values[0] # Simple LR coefficient
print(f"Base LotArea coefficient (Simple LR): {base_coef:.4f}")
print("\nPercentage changes from base model:")
print("-" * 40)

for i, (model_name, coef) in enumerate(zip(short_model_names[1:], coefficients_values[1:])):
    change = ((coef - base_coef) / base_coef) * 100
    print(f"{model_name}: {change:+.2f}%")

# Statistical analysis
print(f"\nStatistical Summary of LotArea Coefficients:")
print(f"Mean: {np.mean(coefficients_values):.4f}")
print(f"Standard Deviation: {np.std(coefficients_values):.4f}")
print(f"Range: {min(coefficients_values):.4f} to {max(coefficients_values):.4f}")

# Compare model performance
print(f"\nModel Performance Summary (Test R²):")
print("-" * 40)

# Get test R² scores for all models
all_test_r2 = []
all_model_names = []

# Original models
for model_name in original_model_names:
    if model_name in summary_results:
        all_test_r2.append(summary_results[model_name]['avg_test_r2'])
        all_model_names.append(model_name.split('(')[0].strip())
```

```

# Mixed models
for model_name in mixed_model_names:
    if model_name in mixed_summary_results:
        all_test_r2.append(mixed_summary_results[model_name]['avg_test_r2'])
        all_model_names.append(model_name.split('(')[0].strip())

for i, (model_name, r2) in enumerate(zip(short_model_names, all_test_r2)):
    print(f"{model_name}: {r2:.4f}")

print("=*80)

```

ANALYSIS OF LOTAREA COEFFICIENT CHANGES

=====

=====

Base LotArea coefficient (Simple LR): 2.1849

Percentage changes from base model:

Model 1: -32.13%

Model 2: -43.45%

Model 3: -69.71%

Model 4: +6.97%

Model 5: +3.99%

Model 6: -49.84%

Statistical Summary of LotArea Coefficients:

Mean: 1.6101

Standard Deviation: 0.6115

Range: 0.6618 to 2.3371

Model Performance Summary (Test R²):

Simple LR: 0.0495

Model 1: 0.1136

Model 2: 0.6509

Model 3: 0.7166

Model 4: 0.0577

Model 5: 0.0995

Model 6: 0.5311

=====

=====